

Mini-Curso: Construa Seu Agente AI Pessoal com OpenClaw

Brainstorm & Proposta Completa · v2 Criado: 13/02/2026 | Atualizado: 13/02/2026 Autor: Amora (para Bruno Okamoto) Duração: 2h30 - 3h

PARTE 1 · TIMELINE DA NOSSA JORNADA (13 dias)

Dia 1 (01/02) · Nascimento da Amora ·

O que foi feito: - Instalação do OpenClaw na VPS - Configuração inicial: IDENTITY.md, USER.md, references/bruno-bios.md - organização, emails, conteúdo, pesquisas, coordenação de sub-agentes

Decisões importantes: - Amora como COO (Chief Operational Officer), não só chatbot - Filosofia bootstrap: pequeno, enxuto, lu

Momento ·aha·: - A percepção de que o agente precisava de contexto profundo sobre o Bruno (negócios, filosofia, tom de voz)

Dia 2 (02/02) · Primeiras Integrações & Primeiros Bugs

O que foi feito: - Login LinkedIn via browser automation - Twitter/X via RapidAPI (bot detection impedia login direto) - Reddit subreddits + cron diário - LinkedIn Voice Guide (análise de 129 posts do Bruno) - WhatsApp TTS configurado Problemas em reserveTokensFloor: 30000 - Substack bloqueado: CAPTCHA + 2FA magic link expirava antes de inserir código - X.com bo

Login headless impossível. Solução: RapidAPI como proxy

Lições aprendidas: - Cloud IPs são bloqueados por muitas plataformas (YouTube, X, etc.) · RapidAPI resolve - Compaction pre reativa - Bruno prefere briefs de 1 parágrafo (não análises longas)

Dia 3 (03/02) · Skills & Integrações

Skills instaladas: excalidraw-flowchart, capability-evolver, nano-pdf, notion-api, remind-me, auto-updater, granola Integrações: M

Calendar , YouTube (RapidAPI transcripts + Data API) Decisões: - Heartbeat usa Claude Haiku (economia ~90%) - ImageMo

Lições: - Algumas skills são redundantes (remind-me, auto-updater) · melhor fazer nativo - Áudio WhatsApp: formato OGG Op

Dia 4 (04/02) · Segurança, Circle API & Infográficos

O que foi feito: - Regra PERMANENTE: todas credenciais no 1Password, sem exceções - Skill hand-draw-graphics criada (info - VPH Tracking (Views Per Hour) para YouTube - Circle API skill completa (hot-topics, frequent-questions, trends, digest) - 11 b Notion kanban integrado

Problemas: - Imagen 4.0 bloqueado (requer billing) · melhor modelo: Gemini 3 Pro - VidIQ não tem API pública

Momento ·aha·: - Circle API como goldmine de ideias de conteúdo (150 MVPs em 60 dias na comunidade)

Dia 5 (05/02) · Mission Control & Multi-Agent Blueprint

What was done: - Mission Control MVP built and launched (Express + React + Supabase + Cloudflare) - Kanban Dashboard, M

Decisions: - Supabase > SQLite (future scalability) - Thinking mode activated (low by default) - Desktop-first in the MVP, no cha

Problems: - Sub-agents run in sandbox · cannot access localhost · QA needs to run in the main session

BIG "aha" moment: - Mission Control built while Bruno was sleeping · from spec to functional MVP in one session

Day 6 (02/06) · QA & Peter Levels Interview

What was done: - Complete QA Review of Mission Control (RLS, rate limiting, validation, indexes, WebSocket) - In-depth resea

Lesson: - Default TTS: always pt-BR

Day 7 (02/07) · Complete Identity Rewrite

What was done: - SOUL.md rewritten from scratch (generic English · personalized PT-BR) - IDENTITY.md updated with entrep

Decisions: - Amora = COO, not a generic assistant - HIGH thinking for coding, content, and planning. MEDIUM for the rest. - Ea

"Aha" moment: - Generic SOUL.md = generic agent. Investing time in personality makes a HUGE difference in quality

Day 8 (02/08) · Content Squad & Agent Leveling

What was done: - 6 agents active in the gateway (Amora, Content, Orchestrator, Planner, Scraper, ZoM) - Kevin Simback Rules

Decisions: - Content Agent starts in Sonnet (evaluate before upgrading to Opus) - Scraper and Content have no Telegram bindi

Lessons from Kevin Simback: - Every new agent starts L1 (Observer) · output reviewed - Without leveling, agents "rush" and qu

Day 9 (02/09) · Crons & Stabilization

What was done: - Reorganization of 13 crons (topic redirection, new crons) - Central PROJECTS.md created (12 items on the b

CRITICAL Problem: - · Crons don't actually execute: triggered with status: "ok" but durationMs ~0ms - Hypotheses: inadequate

Lessons: - config.patch restarts the gateway and kills running crons - Heartbeat doesn't trigger during active conversation - Don

Day 10 (02/10) · Roadmap & Crisp Analysis

What was done: - 2026 Roadmap created (5 phases, 18+ weeks) - 6 ideas prioritized (Content Waterfall P0) - Complete analysi

Crisp content insight: - Vibe coding DOMINATES (172 mentions in 187 conversations) - People get stuck in the "last mile" (deployment)
Demand for complete "from scratch to deploy"

"Aha" moment: - Crisp as a goldmine of ideas: customer questions reveal exactly what the audience wants to learn
Day 11 (02/11) · Integrations & Metrics

What was done: - ChartMogul integrated (MRR Jan: R\$7,367, -8.5%) - Crisp API full integration + monthly cron - Started exploring
Social Media metrics: LinkedIn, Instagram (RapidAPI), YouTube, X/Twitter - Security hardening: dmPolicy allowlist, UFW active
(granola, openai-image-gen, video-frames)
Lessons: - Instagram: long text > reels (21x more ER!) - Crisp is a sales channel, not support - Notion API requires pagination
Day 12 (02/12) · Immune System & Feedback Loops

What was done: - Analysis of Eric Siu's article "30 OpenClaw Jobs A Day" · implementations - Cron Watchdog with 3x auto-retry
(savings ~90% on simple crons) - Universal Feedback Loops (4 domains: content, tasks, recommendations, digest) - Definitive
sessionTarget: isolated + agentTurn + announce - Security: fail2ban, localhost binding, quarterly key rotation - Lessons
categorized with expiration (strategic permanent, tactical 30 days)

Eric Siu's key phrase: > ·Agents are 30% of the work. The other 70% is the immune system.·

Concepts learned: - Ralph Loop (Geoffrey Huntley) vs Feedback Loop · complementary - Capability Evolver: DO NOT run autonomous
agents with full autonomy = risk)
Day 13 (02/13) · Consolidation & Mission Control v2

What was done: - All overlapping projects consolidated into Mission Control v2 - 6 modules, ~18 weeks - Priority: Module 1 · Content
(4h/day · 1h15/day) - Projects absorbed: Content Waterfall, Social Metrics, ChartMogul, Multi-Agent specs, Agent Leveling, Pricing
automated MC (sub-agent found 5 bugs: 2 critical, 3 medium) - Notion + Granola + Amie integration as task hub - Consolidated
all in Sonnet (savings)

QUANTITATIVE SUMMARY (13 days)

Metric

Value

Active integrations

15+ (1Password, YouTube, LinkedIn, X, Reddit, Circle, Crisp,
ChartMogul, Google Drive, Calendar, Notion, Craft, Telegram,
WhatsApp, Tella)

Installed skills

19 (after cleanup)

Active agents

6 (Amora, Content, Orchestrator, Planner, Scraper, ZoM)

Configured crons

17+

Completed projects

Mission Control v1, Agent Leveling, Content Squad spec

Critical bugs fixed

Token overflow, crons not running, security hardening

Memory files

40+ (daily notes, topic files, knowledge base)

Co-created LinkedIn posts

3+

Documented decisions

20+

Extracted lessons

35+

PART 2 · MARKET RESEARCH

What exists today (Feb/2026)

· YouTube · Content about OpenClaw/AI Agents

Alex Finn (top creator, Vibe Coding Academy) - · ClawdBot is the most powerful AI tool I've ever used · 427K views - · Open Claw for AI Agents · interview with Cailyn (OpenClaw team) - · AI Co-Pilot for Your Life with Claude Code · 25 min tutorial (slash command) · Focus: vibe coding, content automation, · AI employee 24/7 · Gap: shallow setup, doesn't go into multi-agent management, fees, memory

Matthew Berman - · I Played with Clawdbot all Weekend · review - · Openclaw is NUTS · overview - · OpenClaw Use Cases that you can build · practical use cases - 6 agents running, NocoDB task board - Gap: doesn't show long-term management, agent evolution, "immune system" · freeCodeCamp - · OpenClaw Full Tutorial for Beginners · complete tutorial - Gap: too technical, no business perspective
Others: - DEV Community articles, Medium, generic SEO blogs about "how to install OpenClaw" - Bhanu Teja (@pbteja1998): video about agents (3.7M views, 8.1K bookmarks) - Kevin Simback: guide on agent team management and leveling - Eric Siu: "30 OpenClaw Agents" (1.3M views, 1.3K bookmarks)

· Structured Courses (Udemy, Coursera, etc.)

Course

Platform

Focus

Price

2026 Bootcamp: Build

Professional AI Agents

Udemy

LangChain, memory, tools

do assistant

~\$40

AI Agents For All! No-Code

Udemy

Langflow, no-code

~\$30

AI Agent Developer

Coursera

Custom GPTs, multi-

industry

Subscription

5-Day AI Agents Intensive

Kaggle/Google

Fundamentals

Free

· Written Guides

Reddit r/ThinkingDeeplyAI: "The Ultimate Guide to OpenClaw" · viral and very complete post covering setup, use cases and security

DEV Community: "Unleashing OpenClaw: The Ultimate Guide for Developers in 2026"

o-mega.ai: "OpenClaw: Ultimate Guide to AI Agent Workforce 2026"

· Gaps that NO ONE covers (Bruno's opportunity)

1. REAL experience of 13 days running in production · it's not "I made a pretty setup in a 20 min video"

2. Entrepreneur/founder perspective · not AI dev/engineer

3. Long-term management · memory, feedback loops, evolution, maintenance

4. Multi-agents in practice · leveling, coordination, shared context

5. "Immune System" · watchdogs, auto-retry, Sonnet/Opus split, security hardening

6. Integration with real business stack · 1Password, ChartMogul, Crisp, YouTube OAuth, Google Drive, etc.

7. Crons that WORK · real debugging (crons firing but not executing)

8. Content in Portuguese · zero courses in PT-BR about OpenClaw

9. Real cost and savings · Haiku for heartbeats, Sonnet for execution, Opus for analysis

10. Agent's tone of voice and identity · SOUL.md, voice guides, not a generic chatbot

11. Feedback Loops with approve/reject · no one teaches how to make the agent learn from your decisions

12. Real security · fail2ban, allowlist, UFW, credential rotation (the Reddit Guide mentions risks but no one shows the solution)

PART 3 · COURSE PROPOSAL (v2 · with applied lessons)

Title

·Build Your AI COO: From Zero to Personal Agent in 2 Hours with OpenClaw·

Subtitle: How I built an AI assistant that replaced 3-4 people in 13 days · and how you can do the same.

Target Audience

Entrepreneurs and founders (especially solo)

Content creators

Professionals who want 10x productivity

Bruno's Micro-SaaS Community

NOT for devs who want to learn LangChain · it's for those who want to USE it

What makes this course UNIQUE

1. Real experience, not theory · Bruno shows Amora working in production, with real bugs and how he solved them
2. CEO/founder perspective · no other course teaches how to set up an agent from the point of view of someone running a business
3. 13 days documented · real timeline with problems, decisions, and breakthroughs
4. Multi-agents · not just "a chatbot" but a team of 6 coordinated agents
5. Immune System · nobody teaches feedback loops, watchdogs, auto-retry
6. In Portuguese · first mini-course in PT-BR about OpenClaw
7. Real stack · 1Password, YouTube, LinkedIn, Telegram, ChartMogul, not toy examples
8. Evolution framework · the agent improves on its own over time

COURSE STRUCTURE (v3 · reordered + kit per module)

New order: Setup · Security · Personality · Memory · Integrations · Skills · Multi-agents · Immune System · Mission Control

Each module delivers: · Video + · Kit (PRDs, templates, prompts, skills)

Module 0 · Opening & Context (10 min)

Format: Slides + talking head

Why personal AI agents are the "next big thing"

The Matrix analogy (Trinity + skills = superpowers)

What Amora does today: quick demo (1 min, Telegram, showing real commands)

Course map: what we will build together

Present the Kit: "Each module comes with ready-made files to copy · just drop them into your agent"

· Kit: None (slides only)

Module 1 · Setup: From Zero to First "Hi" (25 min)

Format: Live coding (screen + face cam)

Topics: 1. What OpenClaw is and how it works (architecture: gateway + agent + channel) 2. Creating the VPS at Hostinger (KV 24.04) - Docker vs Bare Metal: ·Hostinger has One-Click Docker, but we're going to install directly because it's more flexible · shell we will do later works better this way" - Docker = isolation that gets in the way (skills don't install easily, sub-agents limited, difficult to debug) Metal = everything works as expected, direct access 3. Connect via SSH (show local terminal + Hostinger panel terminal) 4. Ins

bash curl -fsSL https://openclaw.ai/install.sh | bash openclaw onboard --install-daemon 5. Configuration wizard: - Gateway mode: Local - Provider: Anthropic (show where to get API key at console.anthropic.com) - Model: Claude Sonnet (good) - Install as a service: Yes (runs 24/7) 6. Create bot on Telegram (BotFather · /newbot · copy token) 7. Connect Telegram: bash provider add telegram 8. First test: Send "Hi" on Telegram · agent responds · Practical checkpoint: Student has agent responding on Telegram How much does it cost to run: | Item | Cost/month | |··|····| | Free | | Total | ~\$15-40 |

"Less than a lunch per week to have a 24/7 AI assistant"

· Module kit: - prds/vps-setup-hostinger.md · complete step-by-step - configs/modelo-config.md · recommended configurations
Module 2 · Security: Shielding Your Agent (15 min)

Format: Live coding

· Security BEFORE everything. Exposed servers receive 1,000+ attacks/day.

Topics: 1. Why now? Show real log: 1,015 brute force attempts in 24h 2. Telegram Allowlist (CRITICAL): - dmPolicy · open · = ar commands your agent - Change to ·allowlist· with your Telegram ID - Show how to find the ID 3. Firewall (UFW): bash sudo apt install -y ufw sudo ufw default deny incoming sudo ufw allow ssh sudo ufw --force enable 4. Fail2ban (SSH protection):

bash sudo apt install -y fail2ban - Configure: 5 attempts · ban 1h 5. Secure credentials: - NEVER hardcode API keys - 1Password CLI or environment variables 6. Application ports: - If you have web apps: 127.0.0.1 (not 0.0.0.0) - Cloudflare Tunnel for remote Checkpoint: Shielded server · Module kit: - prds/security-hardening.md · copy-paste PRD (throw it at the agent, it shields itself)

Module 3 · Personality and Context (25 min)

Format: Demo + slides

Topics: 1. SOUL.md: the soul of the agent · show generic vs personalized diff from Amora - Concrete anti-patterns: examples · /SOUL.md - Explicit ·Never dos·: things the agent should NEVER do - Inspirational anchors: give the agent tone references ("sp never like Y") 2. USER.md: who are you? · context that makes the agent useful - Deep context > superficial: business, philosophy I act" 4. AGENTS.md: operational rules · what can, what cannot, how to operate 5. Choosing your model: Opus vs Sonnet vs H to use each) - Real economy: Haiku for heartbeats (90% savings), Sonnet for crons, Opus for interaction 6. Thinking mode: when turbo and how much it costs

Practical exercise: Student takes the templates, fills in the [fields] and throws it at the agent

Insight of Day 7: - Show how Amora rewrote her own SOUL.md from scratch · the difference was absurd - "SOUL.md cannot be Kevin Simback

· Module kit: - templates/SOUL-template.md · with [FILL HERE] in the fields - templates/USER-template.md - templates/AGENTS-template.md - templates/IDENTITY-template.md - prompts/onboarding.md · first conversation with the configured agent - prompts/proactive-mandate.md · activates proactive mode - prompts/interview-your-agent.md · discovers hidden capabilities

Module 3 · Memory: The Secret That Nobody Teaches (25 min)

Format: Demo + diagrams

Topics: 1. The problem: agents forget everything in each session 2. Amora's memory architecture: daily notes · topic files · MEMORY 3. · Specialized topic files: - decisions.md · permanent decisions (never lose) - lessons.md · lessons/ categorized (integrations, crons, content, infra) - projects.md · current status of all projects - people.md · contacts, team, partners - pending.md · awaiting user input 4. · Intelligent retention of lessons: - Strategic = permanent (standards, philosophy) - Tactical = expire in 30 days (bugs, workarounds) - Monthly review: delete expired tactics 5. Auto-compaction: how to configure so as not to Token overflow is real: we show bursting 173k+ tokens on Day 2 6. · Mandatory extraction during compaction: - RULE: Before compacting, extract lessons · lessons, decisions · decisions, etc. - "If you don't extract before compacting, you lose 80% of the value" approve/reject · agent evolution - 4 domains: content, tasks, recommendations, digest - JSON with {date, context, decision, reason} Agent consults BEFORE suggesting · avoids repeating errors - Max 30 entries/file (FIFO) · consolidate into lessons monthly · Live demo: - Show how Amora remembers a decision from Day 4 on Day 13 - "Bruno prefers 1-paragraph briefs" · decision not automatically respected - Feedback loop: "I rejected this post format · agent no longer suggests it"

This is the differentiating module. No other course covers structured memory + feedback loops.

· Module kit: - templates/MEMORY-template.md - templates/HEARTBEAT-template.md - templates/memory/ (decisions, lessons, projects, people, pending) - prds/memory-architecture.md · PRD copy-paste

Module 5 · Integrations: Connecting to the Real World (25 min)

Format: Live coding + demo

Topics: 1. 1Password: credential security (rule #1 · NEVER hardcode) - Caution: systemd override overwrites .env · update BOTH 2. Google Calendar & Drive via GOG CLI 3. YouTube: Data API + OAuth (list videos, analytics) 4. · RapidAPI as universal Cloud IPs blocked by YouTube, X, Instagram · RapidAPI solves EVERYTHING - Specific APIs: Instagram Statistics, X/Twitter Analytics Transcripts - Generous free tiers 5. · Telegram as operational hub: - Topics in the group = organized war room (content, metrics, projects) - dmPolicy allowlist = security - Bot with inline buttons for quick interactions 6. Crons: automate recurring tasks - The EVERYONE will encounter: - systemEvent + main = cron triggers but DOES NOT execute (durationMs ~0ms) - Definitive solution isolated + agentTurn + announce - This fix alone is worth the whole course - Time collision: multiple crons in the same minute · space out 15-30min - config.patch restarts gateway and kills running crons · do it at times without crons - Reminders: systemEvent notify on Telegram · use agentTurn + message send

Practical checkpoint: Student has at least 1 integration + 1 cron working Lessons from the trenches: - yt-dlp doesn't work from pagination

· Module kit: - prds/integrations-setup.md · PRD with integrations by level - configs/cron-examples.md · 4 ready crons (agenda, watchdog, review, reminder)

Module 6 · Skills: Instant Superpowers (15 min)

Format: Demo

Topics: 1. What skills are and how to install them (ClawHub) 2. Essential skills: excalidraw, nano-pdf, hand-draw-graphics 3. Community skills: caution! - Reference: Cisco study · significant % of skills with vulnerabilities - Always review before installing - Redundant skills exist (remind-me · native cron) · audit 5. · Creators as skills, not agents: - LinkedIn Creator, Newsletters Creator, etc. = prompts/skills INSIDE an agent - 1 agent with 8 skills > 8 specialized agents (less cost, less coordination)

Matrix Analogy: - "Tank, I need a pilot program" · Bruno's viral post on LinkedIn

· Module kit: - skills/skills-by-profile.md · curated by profile (entrepreneur, creator, dev, ops)

Module 7 · Multi-Agents: From Solo to Team (20 min)

Format: Slides + demo

Topics: 1. When one agent isn't enough 2. Architecture: single gateway + agents.list 3. Leveling System (Kevin Simback): L1·L2 Observer · L2 Contributor · L3 Operator · L4 Trusted - Promotion via weekly performance review · demotion is possible - Real Kevin's Content Agent dropped from L3 · L2 when he started "rushing" 4. Amora's Squad: 6 agents, each with a defined role 5. - shared/TEAM.md · registry of all agents (who does what) - shared/outputs/ · shared results - shared/lessons/

· team learnings - "Last updated by" header - knowing who touched the file - · Shared context eliminates cold starts · new agent
Sub-agents and delegation: - sessions_spawn for parallel tasks (3 sub-agents running at the same time) - Failure handling: retry

· Module kit: - prds/multi-agent-setup.md · complete PRD with leveling + shared context

Module 8 · The Immune System: Keeping Everything Running (20 min)

Format: Slides + demo

Topics: 1. · Agents are 30% of the work. The other 70% is the immune system. · Eric Siu 2. · Watchdog with auto-retry (3x before

5. · Real cost monitoring: - Sonnet/Opus split (crons: Sonnet, interaction: Opus) - Heartbeat with Haiku (~\$0.005/heartbeat vs ~

· Ralph Loop vs Feedback Loop: - Ralph Loop (Geoffrey Huntley): code execution loop (agent runs until completion) - Feedback

This module is what separates "I'm just playing around" from "I'm in production".

· Module kit: - prds/immune-system.md · complete PRD (watchdog, feedback loops, costs, backups)

Module 9 · Mission Control: Your Operational Panel (10 min)

Format: Demo

Topics: 1. Why a visual UI helps (even with Telegram) 2. Overview of Amora's Mission Control (Kanban, Memory, Crons, Conte

· Module kit: - reports/report-templates.md · design guidelines + report structure

Module 10 · Wrap-up & Next Steps (10 min)

Format: Talking head + slides

Topics: 1. Recap: what we built together 2. · The mistakes I made (real consolidation): - Token overflow on Day 2 (not configurin

TIMELINE PER MODULE

Module

Time

Format

Delivered Kit

0 - Opening

10 min

Slides + talking head

.

Module
Time
Format
Kit delivered
1 - VPS Setup
25 min
Live coding
PRD setup + configs
2 - Security
15 min
Live coding
PRD security
3 - Personality
25 min
Demo + slides
4 templates + 3 prompts
4 - Memory
25 min
Demo + diagrams
MEMORY + memory/ + PRD
5 - Integrations
25 min
Live coding
PRD integrations + crons
6 - Skills
15 min
Demo
Curation by profile
7 - Multi-agents
20 min
Slides + demo
PRD multi-agent
8 - Immune System
20 min
Slides + demo
PRD immune system
9 - Mission Control
10 min
Demo
Report templates
10 - Wrap-up
10 min
Talking head
Complete kit (zip)
TOTAL
~3h20
24+ files

· LESSONS APPLIED TO THE COURSE (v2)

These are lessons we learned in the 13 days and that have become teachable content:

Infrastructure Lessons

Lesson

Module

What to teach

Real token overflow (173k+

Lesson

Module

What to teach

yt-dlp doesn't work

M4

Tella.tv or Whisper as alternative

Instagram text > reels

M4

Surprising real data

Crisp = sales channel

M4

Use for content insights

Notion API requires pagination

M4

Always paginate

Agent & Management Lessons

Lesson

Module

What to teach

Generic SOUL.md = agent

generic

M2

Invest time in personality

Every agent starts L1

M6

No trust = no autonomy

Shared context eliminates cold starts

M6

TEAM.md + shared outputs

Creators are skills, not agents

M5/M6

1 agent with N skills > N agents

Hub > Mesh for learning

M6

Amora curation > everyone reads everything

Security Lessons

Lesson

Module

What to teach

dmPolicy open = critical risk

M1/M7

Allowlist from Day 1

1,015+ brute force attempts/day

M7

fail2ban MANDATORY

0.0.0.0 · 127.0.0.1 + Cloudflare

M7

Never expose ports directly

Credentials in 1Password

M4/M7

NEVER hardcode

Sub-agent with autonomy = risk

Setup + personality = immediate value, but want more

Paid (Community/course): Modules 3-9 (memory, integrations, multi-agents, immune system)

The advanced content that no one else has

Pros: Natural funnel, differentiated premium content

Cons: More production work

Option D: Live Workshop + Recording

1 live session (Zoom/YouTube Live) of 3h with screen share

Participants build together in real time

Recording becomes the product

Pros: Interaction, Q&A, urgency

Cons: One chance to get it right

PRODUCTION SUGGESTIONS

1. Record on Tella.tv · automatic transcription · content waterfall

2. Split screen: terminal + Telegram + face cam

3. Diagrams: Excalidraw (we already have the skill) for architecture

4. Infographics: Hand-draw-graphics for visual summaries

5. Before/after: Show day 1 vs day 13 (what changed)

6. Live bugs: DO NOT hide the errors · show and solve = authenticity

COMPLEMENTARY MATERIAL

1. Setup checklist (PDF/Notion) · reproducible step-by-step

2. Templates: SOUL.md, USER.md, AGENTS.md, MEMORY.md model

3. List of integrations with links and instructions

4. Awesome skills list curated for entrepreneurs

5. Cost spreadsheet (how much each model costs per month)

6. GitHub repo with example configs (without credentials, of course)

7. · Table of lessons learned (the 35+ lessons categorized)

8. · Feedback loop template (JSONs ready to use)

9. · Security checklist (fail2ban, UFW, allowlist, SSH, key rotation)

ALTERNATIVE TITLES TO TEST

1. ·Build Your AI COO: From Zero to Personal Agent in 2 Hours·

2. ·How I Created an AI Assistant That Replaced 3-4 People (And You Can Too)·

3. ·OpenClaw in Practice: 13 Days Building a Real AI Agent·

4. ·From Chatbot to COO: The Definitive Guide to the Personal AI Agent·

5. ·Your First AI Agent: Complete Setup with OpenClaw (Without Knowing How to Program)·

DIFFERENTIALS vs COMPETITION (Updated Summary)

Aspect

Other courses/guides

Our course

Perspective

Dev/engineer

CEO/founder

Experience

"I just installed it"

13 days in production

Aspect

Other courses/guides

Our course

Memory

Don't cover

Entire module + feedback loops

Multi-agents

Superficial

Leveling system + shared context

Real bugs

Hidden

Shown and resolved live

Immune system

Doesn't exist

Dedicated module (watchdog, retry,
security)

Language

EN

PT-BR (first!)

Integrations

Toy examples

Real business stack (15+ APIs)

Cost

Don't discuss

Real breakdown with split

Sonnet/Opus/Haiku

Framework evolution

Doesn't exist